

Ullas Basavapatna Chandrashekar

+1-(571)237-8792 | ullasbc02@gmail.com | [linkedin.com/in/ullasbc/](https://www.linkedin.com/in/ullasbc/) | github.com/ullasbc02 | Washington DC

TECHNICAL SKILLS

Languages: Java, Python, C++, Rust, SQL, HTML/CSS

Frameworks & Libraries: React, NextJS, SpringBoot, Rails, Flask, LangChain, NLP

Performance & Testing: JMeter, k6, LoadRunner, Robot Framework, API Testing, Selenium, Postman

Other: Grafana, Kibana, Swagger, Dynatrace, Agile, Linux

EXPERIENCE

Systems Engineer

Feb 2022 – Jun 2024

Infosys Ltd

Bengaluru, IN

- Delivered two key **Proofs-of-Concept (POCs)** for the Liberty Global project during transition, including custom Android UI/UX test scripts using Robot Framework and **embedded STB device simulation** and executing k6 performance tests with **10 microservice endpoints** on JioCinema.
- Created Universal JMeter scripts (Stress, Endurance, Load testing) **across 100+ microservices**, significantly streamlining release testing and covering both **Pre-Prod and Production environments**, including night tests for 7 European countries with differing load and **endpoints over 10+ releases**.
- Engineered and executed a large-scale load test for the UEFA league, simulating **1 million synthetic customers** with **Data randomization** to ensure peak capacity readiness of the microservices, reduced **latency by 200ms**, **increasing pods to 40%**.
- Detected and resolved critical performance defects in the prod, including **AWS memory leaks by heap dump analysis**, optimized the recommendation engine **CPU spike** under continuous load by **adding caching**, improved **p99 latency by 25%**.
- Contributed to the migration of end-to-end **legacy on-prem production systems to AWS cloud services** for 3 countries, re-architecting monolithic components to scalable services. Optimized resource allocation and implemented **auto-scaling policies**, resulting in a **reduction of 35% in monthly cloud spend** and manual effort.
- Ranked as a Top Performer during the initial 3-month specialized training in Performance Engineering and low-level design, recognized for technical excellence among a cohort of ~500 engineers.*

Software Engineer

Aug 2021 – Dec 2021

Tranzo Logistics

Bengaluru, IN

- Managed a low-latency logistics monitoring system using Asynchronous JavaScript and a microservices-based REST API layer to ingest and **visualize 5,000+ daily** state transitions.
- Architected a fleet optimization engine using **Gradient Boosted Trees**, achieving an **76% prediction accuracy** for fuel-inefficiency patterns. Engineered features based on **spatial-temporal data** to optimize routing logic, directly reducing carbon intensity by **200+ metric tons per year** through algorithmic fuel-waste mitigation.

NOTABLE PROJECTS

Drifline | *Distributed Systems, Observability, Python, OpenTelemetry*

February 2026

- Automated drift detection platform using **OpenTelemetry and AI-driven** execution-graph diffing to identify microservice dependency shifts and **p95 latency regressions (>2x)** within 1–2 minutes.

Semantic Drift | *NLP, Semantic Drift Detection, Contextual Embeddings, HDBSCAN, Temporal Text Analysis*

March 2026

- Designed a robust, full-cycle system for semantic drift detection across high-velocity social communities, spanning an **NLP pipeline** with **contextual embeddings** (Sentence Transformer, HDBSCAN) and a drift evaluation framework (RoBERTa), and executed a **parallelized inference mechanism for sub-200ms** semantic similarity search latency.

MediMate | *Large Language Model, Natural Language Processing, Semantic Retrieval*

January 2025

- Pioneered **RAG pipeline** utilizing **ChromaDB** for vector storage and semantic retrieval across a **corpus of 30K+ PubMed** citations. Reduced similarity search and inference logic of clinical information lookup **latency by 40%**, while improving preliminary diagnostic **accuracy by 10–15% through LLM-driven reasoning**.

WeCureIt | *SpringBoot, NextJS, PostgreSQL, Firebase, AWS, Ollama, AI Agents*

August 2025

- Integrated AI Agents and Branch and Bound to a full-stack clinic platform to leverage facility–specialty mapping and automate complex scheduling under multi-variable constraints. Scaled REST APIs to sustain **100+ concurrent bookings** with **sub-200ms latency**, leveraging algorithmic pruning to achieve a **70% reduction** in manual scheduling conflicts.

EDUCATION

George Washington University

Master of Science in Computer Science

Washington, DC

May 2026